



Geometry Dash Vision RL Agent

CS231N Milestone 3 • Preliminary Results

Dorian Gulley • James Roberts • Naveen Kannan

Environment Validation

Goal. Train a vision-only agent that maps one rendered game frame to jump/no-jump.



Preliminary Evaluation Setup

Curriculum. Five simple levels: flat, one spike, two spikes, late spike, mixed spikes.

Observation. Single grayscale image frame, shaped `(1, 96, 96)`.


Policy. Tiny CNN trained with REINFORCE.

Baseline. Random binary policy, 30 episodes per level.

Metrics.


- Completion rate

Results: Random vs CNN RL




Finding. Random is surprisingly strong on very easy levels because repeated jump attempts often clear simple spike layouts.

Training Dynamics



Interpretation. REINFORCE occasionally finds successful episodes, but the policy does not stabilize. This matches expectations for sparse/timing-sensitive rewards with one-frame



Analysis and Limitations

Working so far.

- Live visual observations, deterministic rollouts, and reproducible evaluation are working.
- Physics corrections are covered by targeted regression tests.

Current limits.

- CNN + REINFORCE is too noisy for stable jump timing.
- Random remains competitive because the starter levels are very forgiving.

Next Steps

1. Add short frame history or optical-flow-like velocity cues.
2. Replace vanilla REINFORCE with PPO or actor-critic.
3. Tune curriculum so random is weak but solvable behavior remains clear.
4. Add rollout visualizations for learned policies.
5. Report mean $\pm$ std over multiple seeds.

Milestone takeaway. The environment and evaluation harness are working; the learning algorithm is the current bottleneck.